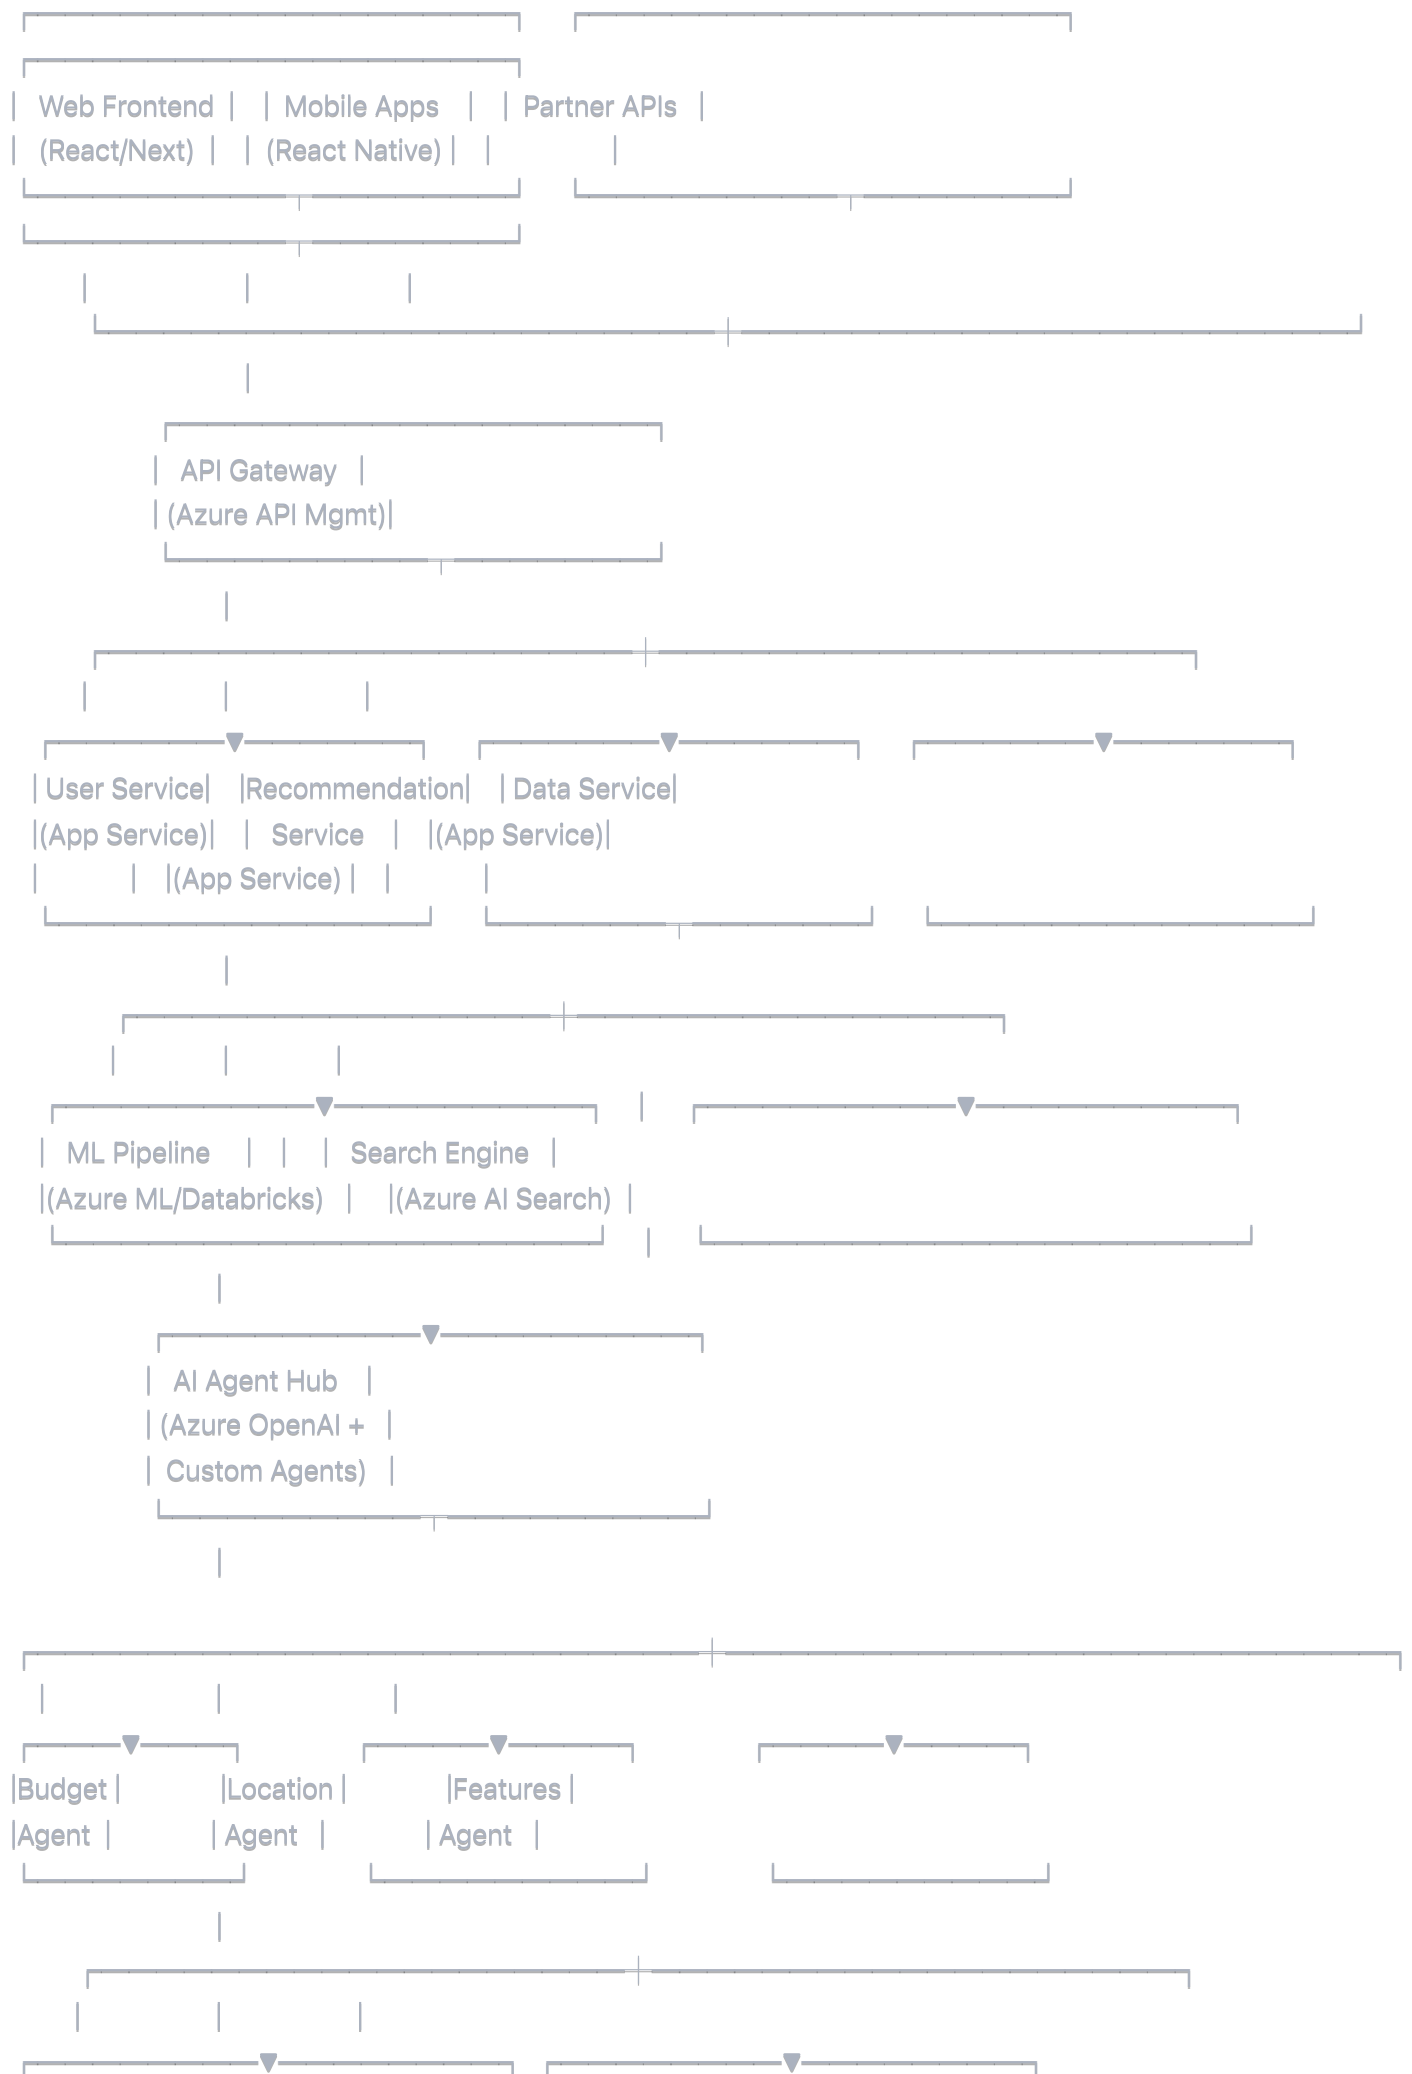



AI Property Recommendation System - Backend Architecture

System Overview

This document outlines the production-ready architecture for an AI-powered property recommendation system, designed to scale on Azure cloud infrastructure.

System Architecture Diagram





Azure Services Mapping

Core Infrastructure

- **Azure App Service:** Host FastAPI backend services
- **Azure API Management:** Gateway, rate limiting, authentication
- **Azure Front Door:** CDN and load balancing
- **Azure Container Registry:** Docker image storage

Data Layer

- **Azure Cosmos DB:** Primary property database (NoSQL)
- **Azure SQL Database:** User profiles and transaction data
- **Azure Blob Storage:** Property images, documents, ML models
- **Azure Data Lake:** Historical data and analytics

AI/ML Layer

- **Azure OpenAI Service:** GPT-4 for explanations and agent coordination
- **Azure Machine Learning:** Custom property scoring models
- **Azure AI Search:** Semantic search and filtering
- **Azure Databricks:** Data processing and model training

Monitoring & Security

- **Azure Application Insights:** Performance monitoring
- **Azure Key Vault:** Secrets and API key management
- **Azure Active Directory B2C:** User authentication
- **Azure Security Center:** Threat detection

API Contracts

1. Property Recommendation Endpoint

POST `/api/v1/recommendations`

Request Schema:

json

```
{
  "user_preferences": {
    "budget_min": 200000,
    "budget_max": 1000000,
    "city": "Denver",
    "min_bedrooms": 2,
    "max_commute_time_minutes": 45,
    "min_school_rating": 7,
    "must_have_features": ["pool", "garage"],
    "lifestyle_preferences": ["modern", "family-friendly"],
    "size_preference": "medium"
  },
  "user_context": {
    "user_id": "user_12345",
    "session_id": "session_67890",
    "previous_searches": []
  },
  "options": {
    "max_results": 3,
    "include_explanations": true,
    "enable_ai_insights": true
  }
}
```

Response Schema:

json

```
{
  "recommendations": [
    {
      "property_id": "prop_98765",
      "address": "123 Example St, Denver, CO",
      "price": 750000,
      "match_score": 92,
      "confidence": 0.87,
      "basic_info": {
        "bedrooms": 3,
        "bathrooms": 2.5,
        "size_sqft": 2200,
        "year_built": 2018,
        "lot_size_sqft": 8000
      },
      "location_info": {
        "city": "Denver",
        "neighborhood": "Capitol Hill",
        "school_rating": 8,
        "commute_time_minutes": 22,
        "walkability_score": 78
      },
      "features": {
        "has_pool": true,
        "garage_spaces": 2,
        "amenities": ["hardwood_floors", "updated_kitchen", "fireplace"]
      },
      "ai_explanation": {
        "summary": "Perfect match for your family with excellent schools and modern amenities",
        "key_reasons": [
          "Within budget with great value at $750k",
          "22-minute commute beats your 45-minute requirement",
          "School rating of 8/10 exceeds your minimum of 7",
          "Modern 2018 construction with desired pool"
        ],
        "potential_concerns": [
          "Slightly smaller lot than neighborhood average"
        ]
      },
      "images": [
        "https://storage.blob.core.windows.net/properties/prop_98765/main.jpg"
      ],
      "agent_insights": {
        "budget_agent_score": 85,
        "location_agent_score": 95,
        "features_agent_score": 90
      }
    }
  ]
}
```

```
    }  
  }  
],  
  "search_metadata": {  
    "total_properties_analyzed": 1247,  
    "filters_applied": 8,  
    "processing_time_ms": 340,  
    "model_version": "v2.1.3"  
  },  
  "ai_summary": "Based on your preferences, I found 3 excellent properties that balance your budget, location, and"   
}
```

2. Property Details Endpoint

GET `/api/v1/properties/{property_id}`

Response Schema:

json

```

{
  "property": {
    "property_id": "prop_98765",
    "detailed_info": {
      "full_address": "123 Example St, Denver, CO 80205",
      "mls_number": "MLS123456",
      "listing_status": "active",
      "days_on_market": 15,
      "price_history": [
        {"date": "2024-01-15", "price": 775000, "event": "price_reduction"},
        {"date": "2024-01-01", "price": 800000, "event": "listed"}
      ]
    },
    "comprehensive_features": {
      "interior": {
        "flooring": ["hardwood", "tile"],
        "kitchen": "updated_2022",
        "appliances": "stainless_steel",
        "heating": "forced_air_gas",
        "cooling": "central_air"
      },
      "exterior": {
        "roof": "composite_shingle_2020",
        "siding": "brick_vinyl",
        "landscaping": "professional",
        "outdoor_features": ["pool", "deck", "garden"]
      }
    },
    "neighborhood_data": {
      "demographics": {
        "median_income": 75000,
        "avg_age": 35,
        "family_percentage": 68
      },
      "nearby_amenities": [
        {"type": "grocery", "name": "Whole Foods", "distance_miles": 0.8},
        {"type": "park", "name": "City Park", "distance_miles": 1.2},
        {"type": "hospital", "name": "Denver Health", "distance_miles": 3.1}
      ]
    }
  }
}

```

3. User Feedback Endpoint

POST `/api/v1/feedback`

Request Schema:

```
json
{
  "user_id": "user_12345",
  "session_id": "session_67890",
  "recommendation_id": "rec_45678",
  "feedback_type": "property_rating",
  "feedback_data": {
    "property_id": "prop_98765",
    "user_rating": 4,
    "feedback_categories": {
      "accuracy": 5,
      "relevance": 4,
      "explanation_quality": 4
    },
    "comments": "Great match, but would prefer larger yard",
    "action_taken": "saved_property"
  }
}
```

AI & Agent Integration

Multi-Agent Architecture

Our system employs specialized AI agents that collaborate to provide comprehensive recommendations:

1. Budget Agent

- **Responsibility:** Financial analysis and affordability assessment
- **Capabilities:**
 - Market value analysis
 - Financing options evaluation
 - Price trend prediction
 - ROI calculations

2. Location Agent

- **Responsibility:** Geographic and commute optimization
- **Capabilities:**
 - Commute time calculation (real-time traffic)
 - Neighborhood analysis

- School district evaluation
- Crime and safety data integration

3. Features Agent

- **Responsibility:** Property amenities and lifestyle matching
- **Capabilities:**
 - Feature importance ranking
 - Lifestyle compatibility scoring
 - Maintenance cost estimation
 - Future upgrade potential

4. Coordinator Agent (Azure OpenAI GPT-4)

- **Responsibility:** Synthesize agent outputs and generate explanations
- **Capabilities:**
 - Natural language explanation generation
 - Conflict resolution between agent recommendations
 - User query interpretation
 - Contextual conversation management

Agent Communication Flow

python

Simplified agent coordination logic

class PropertyRecommendationOrchestrator:

def __init__(self):

 self.budget_agent = BudgetAnalysisAgent()

 self.location_agent = LocationAnalysisAgent()

 self.features_agent = FeaturesAnalysisAgent()

 self.coordinator = GPT4CoordinatorAgent()

async def get_recommendations(self, user_preferences):

Parallel agent execution

 budget_analysis = **await** self.budget_agent.analyze(user_preferences)

 location_analysis = **await** self.location_agent.analyze(user_preferences)

 features_analysis = **await** self.features_agent.analyze(user_preferences)

Coordinator synthesizes results

 recommendations = **await** self.coordinator.synthesize({

 'budget': budget_analysis,

 'location': location_analysis,

 'features': features_analysis,

 'user_context': user_preferences

 })

return recommendations

Edge Case Handling

1. Timeout and Fallback Behavior

python

```
# Timeout handling with circuit breaker pattern
```

```
from circuitbreaker import circuit
```

```
@circuit(failure_threshold=5, recovery_timeout=30)
```

```
async def get_ai_explanation(property_data, user_preferences):
```

```
    try:
```

```
        # Primary AI explanation generation
```

```
        return await openai_service.generate_explanation(
            property_data, user_preferences, timeout=10
        )
```

```
    except TimeoutError:
```

```
        # Fallback to rule-based explanation
```

```
        return generate_rule_based_explanation(property_data, user_preferences)
```

```
    except Exception:
```

```
        # Minimal safe explanation
```

```
        return "This property matches your basic requirements."
```

2. Rate Limiting and Abuse Prevention

```
python
```

```
# Redis-based rate limiting
```

```
from fastapi import HTTPException
```

```
from redis import Redis
```

```
class RateLimiter:
```

```
    def __init__(self, redis_client: Redis):
```

```
        self.redis = redis_client
```

```
    async def check_rate_limit(self, user_id: str, endpoint: str):
```

```
        key = f"rate_limit:{user_id}:{endpoint}"
```

```
        current_requests = await self.redis.get(key)
```

```
        if current_requests and int(current_requests) > 100: # 100 requests per hour
```

```
            raise HTTPException(
                status_code=429,
                detail="Rate limit exceeded. Try again later."
            )
```

```
        await self.redis.incr(key)
```

```
        await self.redis.expire(key, 3600) # 1 hour window
```

3. Data Validation and Security

```
python
```

```
from pydantic import BaseModel, validator
from typing import Optional, List

class UserPreferences(BaseModel):
    budget_min: int
    budget_max: int
    city: Optional[str] = None
    min_bedrooms: int
    max_commute_time_minutes: int

    @validator('budget_min', 'budget_max')
    def validate_budget(cls, v):
        if v < 0 or v > 50_000_000: # Reasonable bounds
            raise ValueError('Budget must be between $0 and $50M')
        return v

    @validator('budget_max')
    def budget_max_greater_than_min(cls, v, values):
        if 'budget_min' in values and v <= values['budget_min']:
            raise ValueError('Maximum budget must be greater than minimum')
        return v

    @validator('city')
    def validate_city(cls, v):
        if v and len(v) > 50:
            raise ValueError('City name too long')
        return v
```

4. Monitoring and Alerting

python

```
# Application Insights integration
from applicationinsights import TelemetryClient
from datetime import datetime

class RecommendationMonitor:
    def __init__(self):
        self.telemetry = TelemetryClient(instrumentation_key="your-key")

    def track_recommendation_request(self, user_id, preferences, results):
        self.telemetry.track_event('RecommendationGenerated', {
            'user_id': user_id,
            'results_count': len(results),
            'processing_time': datetime.now().isoformat(),
            'city_filter': preferences.get('city', 'any')
        })

    def track_error(self, error_type, error_message, user_context):
        self.telemetry.track_exception(
            type=error_type,
            value=error_message,
            properties=user_context
        )
```

Security Best Practices

1. Authentication & Authorization

- Azure AD B2C for user authentication
- JWT tokens with short expiration (15 minutes)
- Refresh token rotation
- Role-based access control (RBAC)

2. Data Protection

- Encryption at rest (Azure Storage encryption)
- Encryption in transit (TLS 1.3)
- PII data masking in logs
- GDPR compliance for EU users

3. API Security

- Input validation and sanitization
- SQL injection prevention

- Rate limiting per user/IP
- CORS policy configuration
- API key management via Azure Key Vault

4. Infrastructure Security

- Network security groups (NSGs)
- Private endpoints for Azure services
- Web Application Firewall (WAF)
- DDoS protection
- Security baselines and compliance scanning

Scalability Considerations

1. Horizontal Scaling

- Auto-scaling App Services based on CPU/memory
- Container orchestration with Azure Kubernetes Service
- Database read replicas for query performance
- Redis cache for session management

2. Performance Optimization

- CDN for static assets (property images)
- Database indexing strategy
- Async processing for heavy operations
- Connection pooling for database connections

3. Cost Optimization

- Reserved instances for predictable workloads
- Spot instances for batch processing
- Storage tiering for historical data
- Function Apps for sporadic tasks

This architecture provides a robust, scalable foundation for the AI-powered property recommendation system while maintaining security, performance, and cost-effectiveness.